

SpaceX_Machine Learning Prediction_Part_5

June 22, 2022

1 Space X Falcon 9 First Stage Landing Prediction

1.1 Assignment: Machine Learning Prediction

Estimated time needed: **60** minutes

Space X advertises Falcon 9 rocket launches on its website with a cost of 62 million dollars; other providers cost upward of 165 million dollars each, much of the savings is because Space X can reuse the first stage. Therefore if we can determine if the first stage will land, we can determine the cost of a launch. This information can be used if an alternate company wants to bid against space X for a rocket launch. In this lab, you will create a machine learning pipeline to predict if the first stage will land given the data from the preceding labs.

Several examples of an unsuccessful landing are shown here:

Most unsuccessful landings are planed. Space X; performs a controlled landing in the oceans.

1.2 Objectives

Perform exploratory Data Analysis and determine Training Labels

- create a column for the class
- Standardize the data
- Split into training data and test data

-Find best Hyperparameter for SVM, Classification Trees and Logistic Regression

- Find the method performs best using test data

1.3 Import Libraries and Define Auxiliary Functions

We will import the following libraries for the lab

```
[61]: import warnings
      warnings.filterwarnings('ignore')
```

```
[62]: # Pandas is a software library written for the Python programming language for
      ↪ data manipulation and analysis.
      import pandas as pd
```

```

# NumPy is a library for the Python programming language, adding support for
↳ large, multi-dimensional arrays and matrices, along with a large collection
↳ of high-level mathematical functions to operate on these arrays
import numpy as np
# Matplotlib is a plotting library for python and pyplot gives us a MatLab like
↳ plotting framework. We will use this in our plotter function to plot data.
import matplotlib.pyplot as plt
#Seaborn is a Python data visualization library based on matplotlib. It
↳ provides a high-level interface for drawing attractive and informative
↳ statistical graphics
import seaborn as sns
# Preprocessing allows us to standardize our data
from sklearn import preprocessing
# Allows us to split our data into training and testing data
from sklearn.model_selection import train_test_split
# Allows us to test parameters of classification algorithms and find the best
↳ one
from sklearn.model_selection import GridSearchCV
# Logistic Regression classification algorithm
from sklearn.linear_model import LogisticRegression
# Support Vector Machine classification algorithm
from sklearn.svm import SVC
# Decision Tree classification algorithm
from sklearn.tree import DecisionTreeClassifier
# K Nearest Neighbors classification algorithm
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics

```

This function is to plot the confusion matrix.

```

[63]: def plot_confusion_matrix(y,y_predict):
    "this function plots the confusion matrix"
    from sklearn.metrics import confusion_matrix

    cm = confusion_matrix(y, y_predict)
    ax= plt.subplot()
    sns.heatmap(cm, annot=True, ax = ax); #annot=True to annotate cells
    ax.set_xlabel('Predicted labels')
    ax.set_ylabel('True labels')
    ax.set_title('Confusion Matrix');
    ax.xaxis.set_ticklabels(['did not land', 'land']); ax.yaxis.
↳ set_ticklabels(['did not land', 'landed'])

```

1.4 Load the dataframe

Load the data

```
[64]: data = pd.read_csv("https://cf-courses-data.s3.us.cloud-object-storage.
↳ appdomain.cloud/IBM-DS0321EN-SkillsNetwork/datasets/dataset_part_2.csv")

# If you were unable to complete the previous lab correctly you can uncomment
↳ and load this csv

# data = pd.read_csv('https://cf-courses-data.s3.us.cloud-object-storage.
↳ appdomain.cloud/IBMDeveloperSkillsNetwork-DS0701EN-SkillsNetwork/api/
↳ dataset_part_2.csv')

data.head()
```

```
[64]:
```

	FlightNumber	Date	BoosterVersion	PayloadMass	Orbit	LaunchSite	\
0	1	2010-06-04	Falcon 9	6104.959412	LEO	CCAFS SLC 40	
1	2	2012-05-22	Falcon 9	525.000000	LEO	CCAFS SLC 40	
2	3	2013-03-01	Falcon 9	677.000000	ISS	CCAFS SLC 40	
3	4	2013-09-29	Falcon 9	500.000000	PO	VAFB SLC 4E	
4	5	2013-12-03	Falcon 9	3170.000000	GTO	CCAFS SLC 40	

	Outcome	Flights	GridFins	Reused	Legs	LandingPad	Block	\
0	None None	1	False	False	False	NaN	1.0	
1	None None	1	False	False	False	NaN	1.0	
2	None None	1	False	False	False	NaN	1.0	
3	False Ocean	1	False	False	False	NaN	1.0	
4	None None	1	False	False	False	NaN	1.0	

	ReusedCount	Serial	Longitude	Latitude	Class
0	0	B0003	-80.577366	28.561857	0
1	0	B0005	-80.577366	28.561857	0
2	0	B0007	-80.577366	28.561857	0
3	0	B1003	-120.610829	34.632093	0
4	0	B1004	-80.577366	28.561857	0

```
[65]: X = pd.read_csv('https://cf-courses-data.s3.us.cloud-object-storage.appdomain.
↳ cloud/IBM-DS0321EN-SkillsNetwork/datasets/dataset_part_3.csv')

# If you were unable to complete the previous lab correctly you can uncomment
↳ and load this csv

X = pd.read_csv('https://cf-courses-data.s3.us.cloud-object-storage.appdomain.
↳ cloud/IBMDeveloperSkillsNetwork-DS0701EN-SkillsNetwork/api/dataset_part_3.
↳ csv')

X.head(100)
```

```

[65]:      FlightNumber  PayloadMass  Flights  Block  ReusedCount  Orbit_ES-L1  \
0           1.0    6104.959412      1.0    1.0           0.0           0.0
1           2.0    525.000000      1.0    1.0           0.0           0.0
2           3.0    677.000000      1.0    1.0           0.0           0.0
3           4.0    500.000000      1.0    1.0           0.0           0.0
4           5.0    3170.000000      1.0    1.0           0.0           0.0
..          ...          ...      ...    ...          ...          ...
85          86.0   15400.000000      2.0    5.0           2.0           0.0
86          87.0   15400.000000      3.0    5.0           2.0           0.0
87          88.0   15400.000000      6.0    5.0           5.0           0.0
88          89.0   15400.000000      3.0    5.0           2.0           0.0
89          90.0    3681.000000      1.0    5.0           0.0           0.0

      Orbit_GEO  Orbit_GTO  Orbit_HEO  Orbit_ISS  ...  Serial_B1058  \
0           0.0         0.0         0.0         0.0  ...           0.0
1           0.0         0.0         0.0         0.0  ...           0.0
2           0.0         0.0         0.0         1.0  ...           0.0
3           0.0         0.0         0.0         0.0  ...           0.0
4           0.0         1.0         0.0         0.0  ...           0.0
..          ...          ...          ...      ...  ...          ...
85          0.0         0.0         0.0         0.0  ...           0.0
86          0.0         0.0         0.0         0.0  ...           1.0
87          0.0         0.0         0.0         0.0  ...           0.0
88          0.0         0.0         0.0         0.0  ...           0.0
89          0.0         0.0         0.0         0.0  ...           0.0

      Serial_B1059  Serial_B1060  Serial_B1062  GridFins_False  GridFins_True  \
0           0.0         0.0         0.0           1.0           0.0
1           0.0         0.0         0.0           1.0           0.0
2           0.0         0.0         0.0           1.0           0.0
3           0.0         0.0         0.0           1.0           0.0
4           0.0         0.0         0.0           1.0           0.0
..          ...          ...          ...          ...          ...
85          0.0         1.0         0.0           0.0           1.0
86          0.0         0.0         0.0           0.0           1.0
87          0.0         0.0         0.0           0.0           1.0
88          0.0         1.0         0.0           0.0           1.0
89          0.0         0.0         1.0           0.0           1.0

      Reused_False  Reused_True  Legs_False  Legs_True
0           1.0         0.0         1.0         0.0
1           1.0         0.0         1.0         0.0
2           1.0         0.0         1.0         0.0
3           1.0         0.0         1.0         0.0
4           1.0         0.0         1.0         0.0
..          ...          ...          ...          ...
85          0.0         1.0         0.0         1.0

```

```
[90 rows x 83 columns]
```

Create a NumPy array from the column Class in data, by applying the method `to_numpy()` then assign it to the variable Y, make sure the output is a Pandas series (only one bracket `df['name of column']`).

-1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -1.05999788e-01,
 -1.50755672e-01, -1.50755672e-01, -1.05999788e-01,
 -1.50755672e-01, -1.50755672e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -1.50755672e-01,
 -1.50755672e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -2.15665546e-01,
 -1.85695338e-01, -2.15665546e-01, -2.67261242e-01,
 -1.05999788e-01, -2.42535625e-01, -1.05999788e-01,
 -2.15665546e-01, -1.85695338e-01, -2.15665546e-01,
 -1.85695338e-01, -1.05999788e-01, 1.87082869e+00,
 -1.87082869e+00, 8.35531692e-01, -8.35531692e-01,
 1.93309133e+00, -1.93309133e+00],
 [-1.67441914e+00, -1.19523159e+00, -6.53912840e-01,
 -1.57589457e+00, -9.73440458e-01, -1.05999788e-01,
 -1.05999788e-01, -6.54653671e-01, -1.05999788e-01,
 -5.51677284e-01, 3.44342023e+00, -1.85695338e-01,
 -3.33333333e-01, -1.05999788e-01, -2.42535625e-01,
 -4.29197538e-01, 7.97724035e-01, -5.68796459e-01,
 -4.10890702e-01, -4.10890702e-01, -1.50755672e-01,
 -7.97724035e-01, -1.50755672e-01, -3.92232270e-01,
 -1.05999788e-01, 9.43398113e+00, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -1.05999788e-01,
 -1.50755672e-01, -1.50755672e-01, -1.05999788e-01,
 -1.50755672e-01, -1.50755672e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -1.50755672e-01,
 -1.50755672e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -2.15665546e-01,
 -1.85695338e-01, -2.15665546e-01, -2.67261242e-01,
 -1.05999788e-01, -2.42535625e-01, -1.05999788e-01,
 -2.15665546e-01, -1.85695338e-01, -2.15665546e-01,
 -1.85695338e-01, -1.05999788e-01, 1.87082869e+00,
 -1.87082869e+00, 8.35531692e-01, -8.35531692e-01,
 1.93309133e+00, -1.93309133e+00],
 [-1.63592675e+00, -1.16267307e+00, -6.53912840e-01,
 -1.57589457e+00, -9.73440458e-01, -1.05999788e-01,
 -1.05999788e-01, -6.54653671e-01, -1.05999788e-01,
 1.81265393e+00, -2.90408935e-01, -1.85695338e-01,
 -3.33333333e-01, -1.05999788e-01, -2.42535625e-01,
 -4.29197538e-01, 7.97724035e-01, -5.68796459e-01,

-4.10890702e-01, -4.10890702e-01, -1.50755672e-01,
 -7.97724035e-01, -1.50755672e-01, -3.92232270e-01,
 -1.05999788e-01, -1.05999788e-01, 9.43398113e+00,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -1.05999788e-01,
 -1.50755672e-01, -1.50755672e-01, -1.05999788e-01,
 -1.50755672e-01, -1.50755672e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -1.50755672e-01,
 -1.50755672e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -2.15665546e-01,
 -1.85695338e-01, -2.15665546e-01, -2.67261242e-01,
 -1.05999788e-01, -2.42535625e-01, -1.05999788e-01,
 -2.15665546e-01, -1.85695338e-01, -2.15665546e-01,
 -1.85695338e-01, -1.05999788e-01, 1.87082869e+00,
 -1.87082869e+00, 8.35531692e-01, -8.35531692e-01,
 1.93309133e+00, -1.93309133e+00],
 [-1.59743435e+00, -1.20058661e+00, -6.53912840e-01,
 -1.57589457e+00, -9.73440458e-01, -1.05999788e-01,
 -1.05999788e-01, -6.54653671e-01, -1.05999788e-01,
 -5.51677284e-01, -2.90408935e-01, -1.85695338e-01,
 3.00000000e+00, -1.05999788e-01, -2.42535625e-01,
 -4.29197538e-01, -1.25356634e+00, -5.68796459e-01,
 2.43373723e+00, -4.10890702e-01, -1.50755672e-01,
 -7.97724035e-01, -1.50755672e-01, -3.92232270e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 9.43398113e+00, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -1.05999788e-01,
 -1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -1.05999788e-01,
 -1.50755672e-01, -1.50755672e-01, -1.05999788e-01,
 -1.50755672e-01, -1.50755672e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -1.50755672e-01,
 -1.50755672e-01, -1.05999788e-01, -1.05999788e-01,
 -1.05999788e-01, -1.50755672e-01, -2.15665546e-01,
 -1.85695338e-01, -2.15665546e-01, -2.67261242e-01,
 -1.05999788e-01, -2.42535625e-01, -1.05999788e-01,
 -2.15665546e-01, -1.85695338e-01, -2.15665546e-01,

```

-1.85695338e-01, -1.05999788e-01, 1.87082869e+00,
-1.87082869e+00, 8.35531692e-01, -8.35531692e-01,
1.93309133e+00, -1.93309133e+00]],
[-1.55894196e+00, -6.28670558e-01, -6.53912840e-01,
-1.57589457e+00, -9.73440458e-01, -1.05999788e-01,
-1.05999788e-01, 1.52752523e+00, -1.05999788e-01,
-5.51677284e-01, -2.90408935e-01, -1.85695338e-01,
-3.33333333e-01, -1.05999788e-01, -2.42535625e-01,
-4.29197538e-01, 7.97724035e-01, -5.68796459e-01,
-4.10890702e-01, -4.10890702e-01, -1.50755672e-01,
-7.97724035e-01, -1.50755672e-01, -3.92232270e-01,
-1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
-1.05999788e-01, 9.43398113e+00, -1.05999788e-01,
-1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
-1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
-1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
-1.05999788e-01, -1.05999788e-01, -1.05999788e-01,
-1.05999788e-01, -1.50755672e-01, -1.05999788e-01,
-1.05999788e-01, -1.50755672e-01, -1.05999788e-01,
-1.50755672e-01, -1.50755672e-01, -1.05999788e-01,
-1.05999788e-01, -1.50755672e-01, -1.50755672e-01,
-1.50755672e-01, -1.05999788e-01, -1.05999788e-01,
-1.05999788e-01, -1.50755672e-01, -2.15665546e-01,
-1.85695338e-01, -2.15665546e-01, -2.67261242e-01,
-1.05999788e-01, -2.42535625e-01, -1.05999788e-01,
-2.15665546e-01, -1.85695338e-01, -2.15665546e-01,
-1.85695338e-01, -1.05999788e-01, 1.87082869e+00,
-1.87082869e+00, 8.35531692e-01, -8.35531692e-01,
1.93309133e+00, -1.93309133e+00]]))

```

We split the data into training and testing data using the function `train_test_split`. The training data is divided into validation data, a second set used for training data; then the models are trained and hyperparameters are selected using the function `GridSearchCV`.

1.7 TASK 3

Use the function `train_test_split` to split the data X and Y into training and test data. Set the parameter `test_size` to 0.2 and `random_state` to 2. The training data and test data should be assigned to the following labels.

X_train, X_test, Y_train, Y_test

```
[69]: X_train, X_test, Y_train, Y_test = train_test_split( X, Y, test_size=0.2,
    ↪ random state=2)
```

we can see we only have 18 test samples.


```
[70]: Y_test.shape
```

```
[70]: (18,)
```

1.8 TASK 4

Create a logistic regression object then create a GridSearchCV object logreg_cv with cv = 10. Fit the object to find the best parameters from the dictionary parameters.

```
[71]: parameters = {"C": [0.01, 0.1, 1], 'penalty': ['l2'], 'solver': ['lbfgs']}# l1 lasso_
      ↪ l2 ridge
      lr=LogisticRegression()
      logreg_cv = GridSearchCV(lr, parameters, cv=10)
      logreg_cv.fit(X, Y)
```

```
[71]: GridSearchCV(cv=10, error_score='raise-deprecating',
      estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
      fit_intercept=True,
      intercept_scaling=1, max_iter=100, multi_class='warn',
      n_jobs=None, penalty='l2', random_state=None, solver='warn',
      tol=0.0001, verbose=0, warm_start=False),
      fit_params=None, iid='warn', n_jobs=None,
      param_grid={'C': [0.01, 0.1, 1], 'penalty': ['l2'], 'solver': ['lbfgs']},
      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
      scoring=None, verbose=0)
```

We output the GridSearchCV object for logistic regression. We display the best parameters using the data attribute best_params_ and the accuracy on the validation data using the data attribute best_score_.

```
[72]: print("tuned hyperparameters :(best parameters) ", logreg_cv.best_params_)
      print("accuracy :", logreg_cv.best_score_)
```

```
tuned hyperparameters :(best parameters) {'C': 0.01, 'penalty': 'l2', 'solver':
'lbfgs'}
accuracy : 0.8222222222222222
```

1.9 TASK 5

Calculate the accuracy on the test data using the method score:

```
[73]: metrics.accuracy_score(Y_test, logreg_cv.predict(X_test))
```


```
[73]: 0.9444444444444444
```

```
[74]: logreg_cv.score(X_test, Y_test)
```

```
[74]: 0.9444444444444444
```

Lets look at the confusion matrix:

```
[75]: yhat=logreg_cv.predict(X_test)
      plot_confusion_matrix(Y_test,yhat)
```



Examining the confusion matrix, we see that logistic regression can distinguish between the different classes. We see that the major problem is false positives.

1.10 TASK 6

Create a support vector machine object then create a GridSearchCV object `svm_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary parameters.

```
[76]: parameters = {'kernel':('linear', 'rbf','poly','rbf', 'sigmoid'),
                    'C': np.logspace(-3, 3, 5),
                    'gamma':np.logspace(-3, 3, 5)}
svm = SVC()
```

```
[77]: svm_cv = GridSearchCV(svm, parameters, cv=10)
      svm_cv.fit(X, Y)
```

```
[77]: GridSearchCV(cv=10, error_score='raise-deprecating',
                  estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
                                decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
                                kernel='rbf', max_iter=-1, probability=False, random_state=None,
```

```

shrinking=True, tol=0.001, verbose=False),
    fit_params=None, iid='warn', n_jobs=None,
    param_grid={'kernel': ('linear', 'rbf', 'poly', 'rbf', 'sigmoid'), 'C':
array([1.00000e-03, 3.16228e-02, 1.00000e+00, 3.16228e+01, 1.00000e+03]),
'gamma': array([1.00000e-03, 3.16228e-02, 1.00000e+00, 3.16228e+01,
1.00000e+03])},
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring=None, verbose=0)

```

```

[78]: print("tuned hpyerparameters :(best parameters) ",svm_cv.best_params_)
      print("accuracy :",svm_cv.best_score_)

```

```

tuned hpyerparameters :(best parameters) {'C': 1.0, 'gamma':
0.03162277660168379, 'kernel': 'sigmoid'}
accuracy : 0.8222222222222222

```

1.11 TASK 7

Calculate the accuracy on the test data using the method score:

```

[79]: metrics.accuracy_score(Y_test, svm_cv.predict(X_test))

```

```

[79]: 0.9444444444444444

```

```

[80]: svm_cv.score(X_test, Y_test)

```

```

[80]: 0.9444444444444444


```

We can plot the confusion matrix

```

[81]: yhat=svm_cv.predict(X_test)
      plot_confusion_matrix(Y_test,yhat)

```



1.12 TASK 8

Create a decision tree classifier object then create a GridSearchCV object `tree_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary parameters.

```
[82]: parameters = {'criterion': ['gini', 'entropy'],
                    'splitter': ['best', 'random'],
                    'max_depth': [2*n for n in range(1,10)],
                    'max_features': ['auto', 'sqrt'],
                    'min_samples_leaf': [1, 2, 4],
                    'min_samples_split': [2, 5, 10]}
```

```
tree = DecisionTreeClassifier()
```

```
[83]: tree_cv = GridSearchCV(tree, parameters, cv=10)
tree_cv.fit(X, Y)
```

```
[83]: GridSearchCV(cv=10, error_score='raise-deprecating',
                  estimator=DecisionTreeClassifier(class_weight=None, criterion='gini',
max_depth=None,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
```

```

        splitter='best'),
        fit_params=None, iid='warn', n_jobs=None,
        param_grid={'criterion': ['gini', 'entropy'], 'splitter': ['best',
'random'], 'max_depth': [2, 4, 6, 8, 10, 12, 14, 16, 18], 'max_features':
['auto', 'sqrt'], 'min_samples_leaf': [1, 2, 4], 'min_samples_split': [2, 5,
10]},
        pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
        scoring=None, verbose=0)

```

```

[84]: print("tuned hpyerparameters :(best parameters) ",tree_cv.best_params_)
      print("accuracy :",tree_cv.best_score_)

```

```

tuned hpyerparameters :(best parameters) {'criterion': 'gini', 'max_depth': 4,
'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 2,
'splitter': 'best'}
accuracy : 0.8888888888888888

```

1.13 TASK 9

Calculate the accuracy of tree_cv on the test data using the method score:

```

[85]: metrics.accuracy_score(Y_test, tree_cv.predict(X_test))

```

```

[85]: 0.9444444444444444

```

```

[86]: tree_cv.score(X_test, Y_test)

```

```

[86]: 0.9444444444444444


```

We can plot the confusion matrix

```

[87]: yhat = tree_cv.predict(X_test)
      plot_confusion_matrix(Y_test,yhat)

```



1.14 TASK 10

Create a k nearest neighbors object then create a GridSearchCV object knn_cv with cv = 10. Fit the object to find the best parameters from the dictionary parameters.

```
[88]: parameters = {'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
                    'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],
                    'p': [1,2]}
```

```
KNN = KNeighborsClassifier()
```

```
[89]: knn_cv = GridSearchCV(KNN, parameters, cv=10)
knn_cv.fit(X, Y)
```

```
[89]: GridSearchCV(cv=10, error_score='raise-deprecating',
                  estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,
                  metric='minkowski',
                  metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                  weights='uniform'),
                  fit_params=None, iid='warn', n_jobs=None,
                  param_grid={'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10], 'algorithm':
                  ['auto', 'ball_tree', 'kd_tree', 'brute'], 'p': [1, 2]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                  scoring=None, verbose=0)
```

```
[90]: print("tuned hpyerparameters :(best parameters) ",knn_cv.best_params_)  
      print("accuracy :",knn_cv.best_score_)
```

```
tuned hpyerparameters :(best parameters) {'algorithm': 'auto', 'n_neighbors':  
5, 'p': 1}  
accuracy : 0.8444444444444444
```

1.15 TASK 11


Calculate the accuracy of tree_cv on the test data using the method score:

```
[91]: knn_cv.score(X_test, Y_test)
```

```
[91]: 0.9444444444444444
```

We can plot the confusion matrix

```
[92]: yhat = knn_cv.predict(X_test)  
      plot_confusion_matrix(Y_test,yhat)
```



1.16 TASK 12

Find the method performs best:

```
[94]: # Compare the best_scores of these tests:
      # The decision tree had the highest best accuracy of 0.8888!
```

The test with the highest accuracy was the decision tree test!

1.17 Authors

Joseph Santarcangelo has a PhD in Electrical Engineering, his research focused on using machine learning, signal processing, and computer vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

1.18 Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2021-08-31	1.1	Lakshmi Holla	Modified markdown
2020-09-20	1.0	Joseph	Modified Multiple Areas

Copyright © 2020 IBM Corporation. All rights reserved.